

MIT ES.S20 Lecture 1: Game Representation

Lecture Outline

1. What do we mean by toys and games?
2. Why do we care about this?
3. How do we represent games mathematically?
4. How does searching on game trees work?
5. What happens when you can't see the end?
6. First Game Demonstration

Toys and Games

- Examples of toys

Geometric Puzzles, Rubik's Cube, Sudoku, Kakuro

- Examples of games

Chess, Checkers, Go, Backgammon, Dominion, Monopoly

- Strict Definitions

To define a game, you need n adversarial players, an exact set of rules, and an impartial judge. You don't need turns per se – this is the boundary between combinatorial (turn- based) and economic (simultaneous or continuous) games.

In toys the rules are generally physical and there is usually only one player.

- Recreational Mathematics

Recreational mathematics is what happens when you take games as inspiration for mathematical problems:

“All mathematics is a recreation” - Conway

The Point of this Course

- Games as Problems

The initial approach is to see a game as nothing more than a specific problem to solve: this is somewhat the motivation of the pure-mathematical approach. It brings up questions like “Why is Go harder than Chess?”; “What happens when this choice is random?”; and “Is this even fair?”

- Games as Cognition

The way people play games, however, is much less precise: it’s very much in terms of visual and kinetic processing. In this way, if you’re trying to think about cognition, specifically machine cognition, a game with tight rules is a good place to start.

- Games as Mental Models

In a strict mathematical context, games are just complex algebraic objects; so they can be used to think about math by proxy. Conway’s famous example is of mapping the real numbers onto the “Surreal Numbers” formed by a simple game; but we will think mostly about the Rubik’s Cube.

- Games as Economics

If you generalize it in that way, much of modern society takes the form of a game: stock markets are actually just a set of rules in an environment, enforced by the government. Game theory is huge business in economics.

- Games as Security

What happens when you want to implement your own game? Well, you can do this directly by designing them in code or legalese in order to make a desired economy, as in MMOs.

A very interesting application is applying games to cryptography when your rules can be enforced by cryptography, as in cryptocurrencies.

Representation of Games

- Game States

Let's take simple, turn-based, complete knowledge games for now; and specifically, we'll solve Tic-Tac-Toe. There are many other ways to make game solutions, but this is probably the simplest.

Whatever is on the board we call a “state” and a move by either player is an edge between these states (a ply).

- Game Trees

We can construct a game tree by connecting the state nodes by edges describing the possible moves. In games like tic-tac-toe where the players alternate turns, each depth level alternates players.




Figure 1: Game Tree Example

We call the number of moves that can be made from any given game state the branching factor

Game Searching

- Minimax Search

If we want to maximize our performance, then we have to search for the maximum performance on our move edges and the minimum on our opponent's. With this in mind, if we want to search for the move that will maximize our performance, we can recurse with the following two functions (in python pseudocode):

```
def max-value(state,depth):
    if (depth == 0):
        return state.value
    v = -float('inf')
    for each s in state.children:
        v = max(v,min-value(s,depth-1))
    return v

def min-value(state,depth):
    if (depth == 0):
        return state.value
    v = -float('inf')
    for each s in state.children:
        v = min(v,max-value(s,depth-1))
    return v
```

In this code, v is the valuation of the position. Take for example 1 as a win, 0 as a draw, and -1 as a loss.

It's instructive to see why we can also use Negamax search, a slight variation which uses the property that $-\min(-a, -b) = \max(a, b)$ to combine the two into one function.

```
def negamax(state,depth):
    if (depth == 0):
        return state.value
    v = -float('inf')
    for each s in state.children:
        v = max(v,-negamax(s,depth-1))
    return v
```

- Iterative Deepening

Note that Minimax/Negamax goes all the way to the bottom of the tree immediately, which will be impossible to do in games like Chess. So, what we can do instead is to only go to a certain depth in an iterative process.

Note, however, that this is difficult because you have to know in advance how far you're going to go and you have to have an evaluation function. We'll talk about how to optimize this sort of procedure when we talk about Chess.

- Heuristic Search

Maybe raw negamax will solve games like tic-tac-toe, but in games with many moves, we need to be able to tell which part of the tree (i.e. which states) are better than others. We call this our evaluation function or heuristic; and if we have a good one, we can just add it to our valuation. If we have a specific goal in mind, this is called A*.

```
def heuristic_search(state, max_depth):
    if (max_depth == 0 or state.is_end()):
        return evaluate(state)

    v = -float('inf')
    for each s in state.children:
        v = max(v, -heuristic_search(s, max_depth-1))
    return v
```

If we accept feedback from the success and failure of our heuristic function, we can learn a proper heuristic. We'll see this later when we talk about artificial intelligence

- Monte-Carlo Methods

Once we're in the context of heuristics, we can often say (as is the case in games like Go or many probabilistic games like Poker) that there are just too many branches and a lot of them look the same. When this happens, randomized strategies like Monte-Carlo make a lot of sense:

The assumption is that simulating play at a random node will be representative and we can get a rough idea of the probability of a good outcome by trying random continuations. We'll talk more about this when we deal with randomized games.

- Proof Searching

Another way of thinking about the search methods we've been describing is in the context of mathematical proofs: if we let a provable win for us be 1 and a loss be 0, we can say that we can win iff $\text{OR}(\text{all of our moves})$. Similarly, if it is our opponent's turn, they can win iff $\text{AND}(\text{all of their moves})$ is 0.

Clearly, this turns the problem of game searching into a giant boolean circuit and is particularly relevant when we start talking about possibilities of winning. If we want to attach probabilities of winning to nodes instead of 1 and 0, we can use fuzzy logic, but this ends up being just the same as heuristic search.

Demonstration

- Strategic Positioning Game
- Loopy
- Playing chess?